


```
Suggested code may be subject to a license | www.kaggle.com/tsheposono/beyond-linear-regression | px-wing.hatenablog.com/entry/2022/03/25/080758 | www.analyticsvidhya.com/blog/2019/08/decoding
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
from google.colab import files
uploaded = files.upload()
```




Choose Files Algerian\_forest\_fires\_cleaned (1).csv


- Algerian\_forest\_fires\_cleaned (1).csv(text/csv) - 15094 bytes, last modified: 6/4/2024 - 100% done

Saving Algerian\_forest\_fires\_cleaned (1).csv to Algerian\_forest\_fires\_cleaned (1) (2).csv

```
import pandas as pd
df = pd.read_csv('Algerian_forest_fires_cleaned (1).csv')
df.head()
```



	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	not fire	0
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	not fire	0
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire	0
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	not fire	0
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	not fire	0



Next steps:

Generate code with df

 View recommended plots

```

import pandas as pd
import numpy as np
from google.colab import files
import matplotlib.pyplot as plt

# Inspect the first few rows of the dataframe
print(df.head())

# Check for missing values
print(df.isnull().sum())

# If there are categorical columns with missing values, fill them with the mode
# df['categorical_column'].fillna(df['categorical_column'].mode()[0], inplace=True)

# Check for data types
print(df.dtypes)

# Convert data types if necessary (Example: day, month, year should be integers)
df['day'] = df['day'].astype(int)
df['month'] = df['month'].astype(int)
df['year'] = df['year'].astype(int)

# Handling outliers (Example: Removing outliers beyond 3 standard deviations)
numeric_columns = ['Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI']
for col in numeric_columns:
    mean = df[col].mean()
    std = df[col].std()
    df = df[(df[col] > mean - 3*std) & (df[col] < mean + 3*std)]

# Feature Selection
# Assuming 'Classes' is the target variable and 'Region' might be categorical
selected_features = ['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Region']
df = df[selected_features + ['Classes']]

# Creating a new feature 'season' based on 'month'
def get_season(month):
    if month in [12, 1, 2]:
        return 'Winter'
    elif month in [3, 4, 5]:
        return 'Spring'
    elif month in [6, 7, 8]:
        return 'Summer'
    else:
        return 'Autumn'

df['season'] = df['month'].apply(get_season)

# Convert categorical features to numerical (One-Hot Encoding)
df = pd.get_dummies(df, columns=['season', 'Region'], drop_first=True)

# Final dataset
print(df.head())

# Save the cleaned and preprocessed dataset to a new CSV file
df.to_csv('cleaned_preprocessed_data.csv', index=False)

# Download the cleaned and preprocessed file
files.download('cleaned_preprocessed_data.csv')

```

```

day month year Temperature RH Ws Rain FFMC DMC DC ISI BUI \
0 1 6 2012 29 57 18 0.0 65.7 3.4 7.6 1.3 3.4
1 2 6 2012 29 61 13 1.3 64.4 4.1 7.6 1.0 3.9
2 3 6 2012 26 82 22 13.1 47.1 2.5 7.1 0.3 2.7
3 4 6 2012 25 89 13 2.5 28.6 1.3 6.9 0.0 1.7
4 5 6 2012 27 77 16 0.0 64.8 3.0 14.2 1.2 3.9

FWI Classes Region
0 0.5 not fire 0
1 0.4 not fire 0
2 0.1 not fire 0
3 0.0 not fire 0
4 0.5 not fire 0
day 0
month 0
year 0
Temperature 0
RH 0
Ws 0
Rain 0
FFMC 0
DMC 0
DC 0
ISI 0
BUI 0
FWI 0
Classes 0
Region 0
dtype: int64
day int64
month int64
year int64
Temperature int64
RH int64
Ws int64
Rain float64
FFMC float64
DMC float64
DC float64
ISI float64
BUI float64
FWI float64
Classes object
Region int64
dtype: object
day month year Temperature RH Ws Rain FFMC DMC DC ISI BUI \
0 1 6 2012 29 57 18 0.0 65.7 3.4 7.6 1.3 3.4
1 2 6 2012 29 61 13 1.3 64.4 4.1 7.6 1.0 3.9
4 5 6 2012 27 77 16 0.0 64.8 3.0 14.2 1.2 3.9
5 6 6 2012 31 67 14 0.0 82.6 5.8 22.2 3.1 7.0
6 7 6 2012 33 54 13 0.0 88.2 9.9 30.5 6.4 10.9

FWI Classes season_Summer Region_1
0 0.5 not fire True False
1 0.4 not fire True False
4 0.5 not fire True False

```

```
import matplotlib.pyplot as plt
import seaborn as sns



# considering 'df' is your DataFrame containing the dataset

# Plotting Temperature distribution
plt.figure(figsize=(8, 6))
sns.histplot(df['Temperature'], bins=20, kde=True, color='skyblue')
plt.title('Temperature Distribution')
plt.xlabel('Temperature')
plt.ylabel('Frequency')
plt.show()

# Plotting Relative Humidity (RH) distribution
plt.figure(figsize=(8, 6))
sns.histplot(df['RH'], bins=20, kde=True, color='salmon')
plt.title('Relative Humidity (RH) Distribution')
plt.xlabel('RH')
plt.ylabel('Frequency')
plt.show()

# Plotting Wind Speed (Ws) distribution
plt.figure(figsize=(8, 6))
sns.histplot(df['Ws'], bins=20, kde=True, color='lightgreen')
plt.title('Wind Speed (Ws) Distribution')
plt.xlabel('Ws')
plt.ylabel('Frequency')
plt.show()

# Plotting Rain distribution
plt.figure(figsize=(8, 6))
sns.histplot(df['Rain'], bins=20, kde=True, color='gold')
plt.title('Rain Distribution')
plt.xlabel('Rain')
plt.ylabel('Frequency')
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder

# Encode categorical target variable
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)

# One-hot encode the target variable
from sklearn.preprocessing import OneHotEncoder
onehot_encoder = OneHotEncoder(sparse=False)
y_train_encoded = y_train_encoded.reshape(len(y_train_encoded), 1)
y_train_onehot = onehot_encoder.fit_transform(y_train_encoded)

# Linear Regression
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train_onehot)

# Polynomial Regression
poly_reg = LinearRegression()
poly_reg.fit(X_train_poly, y_train_onehot)
```

→ /usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/\_encoders.py:868: FutureWarning: `LabelEncoder` will accept no keyword arguments in the future. Please use `LabelEncoder.fit_transform` instead.

▼ LinearRegression  
LinearRegression()

```


from sklearn.preprocessing import LabelEncoder

# Encode categorical target variable
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)

# Lasso Regression
lasso_reg = Lasso(alpha=0.1)
lasso_reg.fit(X_train_scaled, y_train_encoded)

# Ridge Regression
ridge_reg = Ridge(alpha=0.1)
ridge_reg.fit(X_train_scaled, y_train_encoded)

```

 **Ridge**  
Ridge(alpha=0.1)

```

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import Lasso, Ridge
from sklearn.metrics import mean_squared_error

# Split the data into features (X) and target variable (y)
X = df[['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI']]
y = df['Classes']

# Encode the target variable (assuming non-numeric)
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=42)

# Create a pipeline with scaling and a regression model (Lasso or Ridge)
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('regression', Lasso()) # Change to Ridge if desired
])

# Define a parameter grid for hyperparameter tuning
param_grid = {
    'regression__alpha': [0.1, 0.5, 1.0] # Adjust alpha values as needed
}


# Perform GridSearchCV for hyperparameter tuning
grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

# Get the best model from the grid search
best_model = grid_search.best_estimator_

# Evaluate the best model on the test set
y_pred = best_model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)

print("Mean Squared Error:", mse)

```

 Mean Squared Error: 0.10700365666244754

Effective use of cross validation : Cross-validation is like testing how good your cooking skills are by making different dishes multiple times. Instead of just making one dish and hoping it turns out well, you make several dishes using different ingredients and recipes. This way, you get a better idea of how well you cook overall.

Similarly, in machine learning, cross-validation is like testing how good a model is by training it on different parts of the data multiple times. Instead of just training once and hoping for the best, you train the model multiple times on different parts of the data and see how well it performs on average.

This technique helps in choosing the best model, tweaking its settings (like adjusting the heat when cooking), and making sure it doesn't just work well on the data it was trained on, but also on new data it hasn't seen before. Overall, cross-validation helps in making sure your model is